
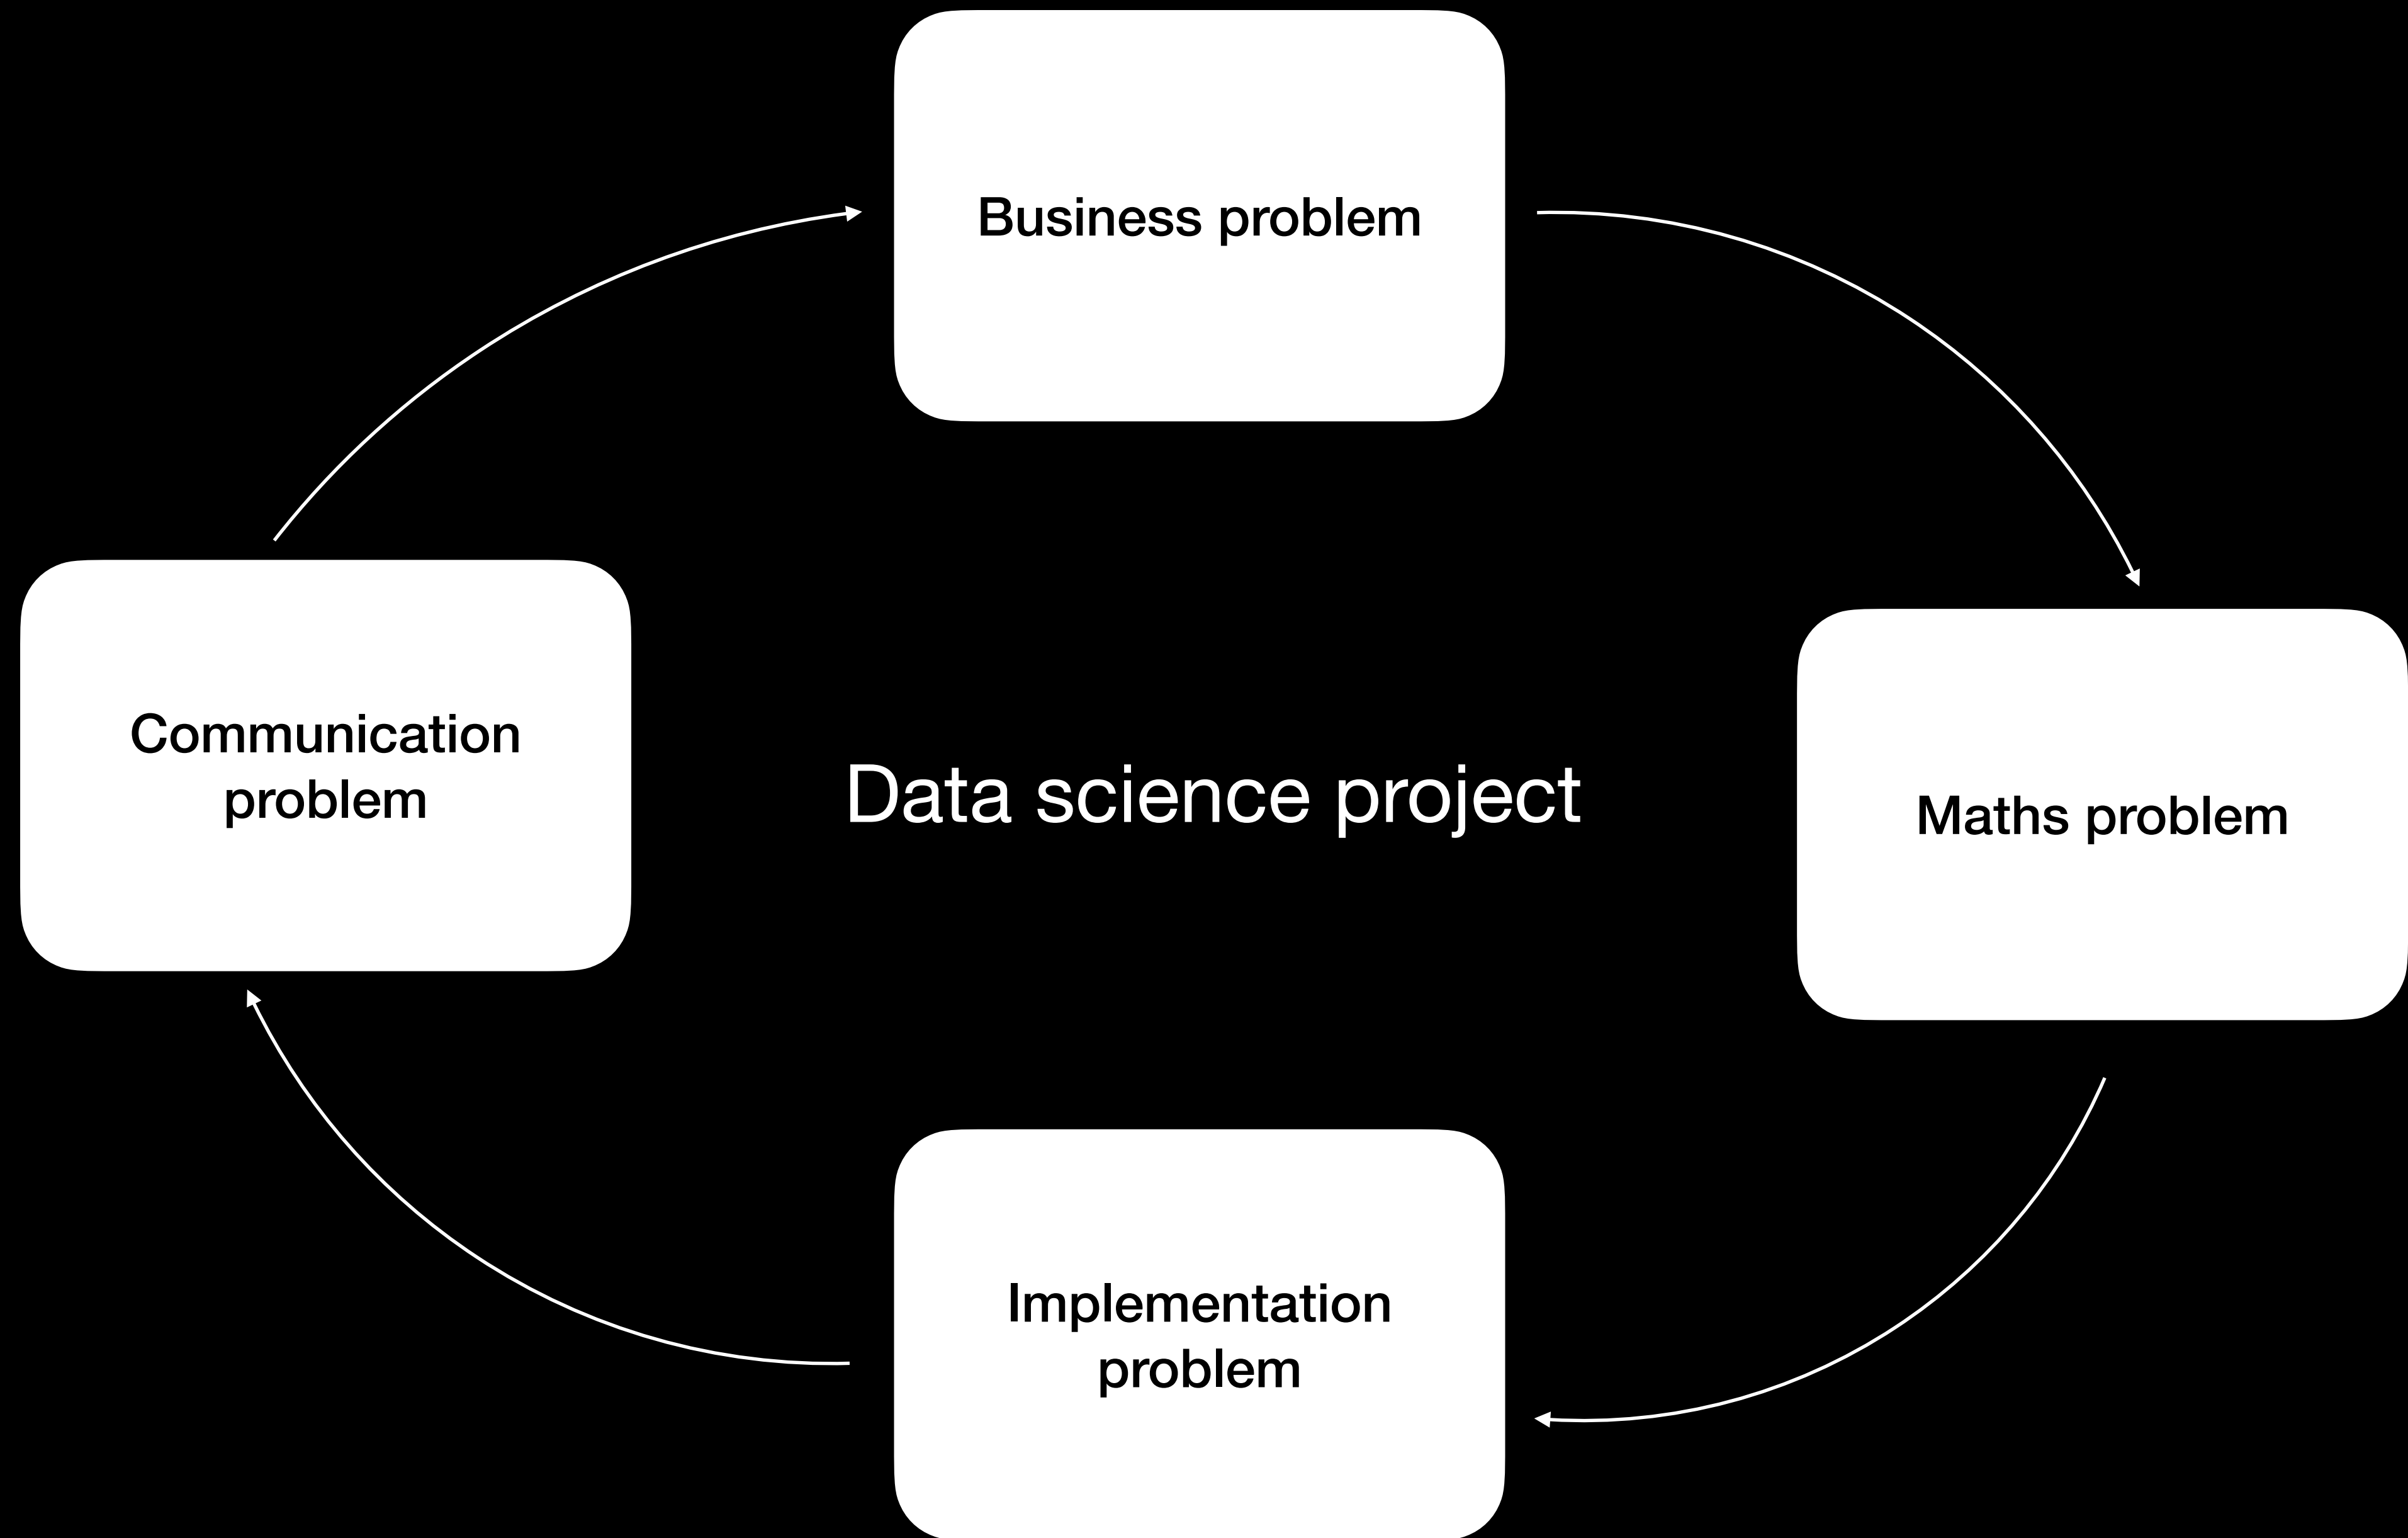


# **MADS - Deployment**

**It works on my machine**

**Raoul Grouls, 2 December 2024**





# Docker for Data Scientists: Getting Started

## Why containerization matters for ML/AI projects

Promise of reproducibility across different machines

- Same environment
- Same dependencies
- Same behavior
- Isolation of dependencies from your personal workflow

Easy sharing of complete project setups

- Less "it works on my machine"
- Quick onboarding of new team members (if they have enough RAM for Docker...)

Simplified deployment process

- From local to production with minimal changes



# Docker for Data Scientists: Getting Started

Why containerization matters for ML/AI projects

Easy sharing of complete project setups

- Less "it works on my machine"
- Quick onboarding of new team members (if they have enough RAM for Docker...)

Simplified deployment process

- From local to production with minimal changes





# The "Works on My Machine" Paradox

Docker doesn't completely solve architecture differences:

- CPU Architecture Challenges
  - x86\_64 (Intel/AMD) vs ARM64 (M1/M2 Macs)
  - Some packages aren't available for all architectures
  - Performance can vary significantly

- Common Issues
  - TensorFlow/PyTorch optimizations differ
  - C++ extensions might not compile
  - Binary dependencies might be architecture-specific
- Possible Solutions
  - Use `--platform` flag when building
  - Build multi-architecture images
  - Document architecture requirements clearly

—\\_ (ツ) \\_—  
**IT WORKS**  
on my machine

# Cross-Platform Docker: Using Platform Flags

Force x86 build on M1 Mac

- `docker build --platform linux/amd64 -t mymodel:latest .`

Force ARM build on any machine

- `docker build --platform linux/arm64 -t mymodel:latest .`

Run for a speccific platform

- `docker run --platform linux/amd64 my-image`

# Docker Fundamentals: Key Concepts

## Core Components:

- **Dockerfile** : The build instructions for creating an **Image**. Similar to a recipe with step-by-step instructions
- **Image** : The blueprint/template for **Containers**. Like a snapshot that can create identical containers
- **Container** : A running instance with your application and dependencies. Think of it as a lightweight VM that shares the host OS kernel
- **Registry** : A repository for storing and sharing images. Can be public (Docker Hub) or private



# Docker Fundamentals: Key Concepts

**Dockerfile** -> builds -> **Image** -> runs -> **Container**

# Understanding the Dockerfile

## FROM

The command FROM:

Docker checks your local machine:

- Looks in local image cache
- If found, uses cached image
- If not found, needs to download

If not found locally:

- Connects to registry (default: Docker Hub)
- Downloads the image in layers
- Saves in local cache for future use

```
1 FROM python:3.12-slim
2
3 WORKDIR /app
4
5 COPY requirements.txt requirements.txt
6 COPY src src
7
8 RUN pip install --no-cache-dir --upgrade -r requirements.txt
9
10 EXPOSE 8000
11
12 ENTRYPOINT ["uvicorn", "src.main:app", "--host", "0.0.0.0", "--port", "8000"]
13
```

# Understanding the Dockerfile

## FROM python:3.12-slim

There are multiple python images, you can find them on [https://hub.docker.com/\\_/python/tags](https://hub.docker.com/_/python/tags)

- -alpine
  - Ultra minimal (50MB)
  - Based on alpine linux
  - Security focused
  - Harder to debug
- -slim
  - Minimal image (130MB)
  - Only essentials
- Good for simple webapps like fastapi when size matters but alpine is too minimal
- Python:3.x
  - Full image (920MB)
  - Includes common system libraries
  - Good for complex dependencies
  - Larger deployment size
- -slim-bullseye / -slim-bookworm
  - Bullseye = Debian 11
  - Bookworm = Debian 12

# Understanding the Dockerfile

## WORKDIR

WORKDIR /app

- Sets the working directory
- All subsequent commands run here
- Good practice for organization

```
1 FROM python:3.12-slim
2
3 WORKDIR /app
4
5 COPY requirements.txt requirements.txt
6 COPY src src
7
8 RUN pip install --no-cache-dir --upgrade -r requirements.txt
9
10 EXPOSE 8000
11
12 ENTRYPOINT ["uvicorn", "src.main:app", "--host", "0.0.0.0", "--port", "8000"]
13
```

# Understanding the Dockerfile

## COPY

`COPY requirements.txt .`

- Copies files from host to container
- Best practice: Copy requirements first
- Helps with layer caching

```
1 FROM python:3.12-slim
2
3 WORKDIR /app
4
5 COPY requirements.txt requirements.txt
6 COPY src src
7
8 RUN pip install --no-cache-dir --upgrade -r requirements.txt
9
10 EXPOSE 8000
11
12 ENTRYPOINT ["uvicorn", "src.main:app", "--host", "0.0.0.0", "--port", "8000"]
13
```

# Understanding the Dockerfile

## RUN

`RUN pip install -r requirements.txt`

- Executes commands during build
- Creates a new layer in the image
- Use `&&` for multiple commands to reduce layers

```
1 FROM python:3.12-slim
2
3 WORKDIR /app
4
5 COPY requirements.txt requirements.txt
6 COPY src src
7
8 RUN pip install --no-cache-dir --upgrade -r requirements.txt
9
10 EXPOSE 8000
11
12 ENTRYPOINT ["uvicorn", "src.main:app", "--host", "0.0.0.0", "--port", "8000"]
13
```



# CMD vs ENTRYPOINT

## Making Your Container Executable

### CMD:

- Can be easily overridden from command line
- If you provide runtime arguments, entire CMD is replaced
- Good for default behavior that might change

### ENTRYPOINT:

- Runtime arguments are added to ENTRYPOINT arguments
- Harder to override (needs `--entrypoint` flag)
- Good for containers that always run the same command

# CMD vs ENTRYPOINT

## Making Your Container Executable

```
# Dockerfile
```

```
CMD ["python", "app.py"]
```

```
# Runtime - completely replaces CMD
```

```
docker run myimage echo "hello" # runs: echo "hello"
```

# CMD vs ENTRYPOINT

## Making Your Container Executable

```
# Dockerfile
```

```
ENTRYPOINT ["python"]
```

```
# Runtime - adds to ENTRYPOINT
```

```
docker run myimage app.py # runs: python app.py
```

```
docker run myimage -V      # runs: python -V
```

# CMD vs ENTRYPOINT

## Making Your Container Executable

Best Practices:

- Use ENTRYPOINT for "main executable" that should always run
- Use CMD for arguments that might change
- Always use JSON array format ["command", "arg"]

```
ENTRYPOINT ["uvicorn"]
```

```
CMD ["main:app", "--host", "0.0.0.0", "--port", "8000"]
```

# Interactive Docker

## Great for debugging

# Only overrides the CMD

```
docker run -it my-image /bin/bash
```

# for Alpine-based images

```
docker run -it my-image /bin/sh
```

# Ignores both ENTRYPOINT and CMD

```
docker run -it --entrypoint /bin/bash my-image
```

# Interactive Docker

Great for debugging

```
# Check installed Python packages  
pip list
```

```
# Test imports  
python  
>>> import numpy
```

```
# Check environment variables  
env
```

```
# Check file system  
ls -la  
cat requirements.txt
```



# Makefile

## Using Makefile to automate the docker commands

```
1 build:
2     @echo "building docker image"
3     docker build -t fastapi-app:latest .
4
5 run:
6     docker run -p 8000:8000 fastapi-app:latest
7     @echo "running on http://localhost:8000"
8
9 interactive:
10    @echo "Entering container interactively"
11    docker run -it --entrypoint /bin/bash fastapi-app:latest
12
```